


Service time is censored: it counts only the tasks that finished

Both panels are the same five planners on corridors, in the same order.

Compare service times only between planners of comparable throughput; on its own, the metric rewards giving up.

tasks delivered per 1000 timesteps
higher is better



mean service time, timesteps
lower looks better

