

Six hypotheses, each scored on the quantity it actually claims to move

Treatment against the control that isolates its mechanism. A verdict needs $p < 0.05$ on the hypothesis's own metric; each cell prints treatment vs control.

H1 · aisle_throughput_per_1000 higher is better	85.9 vs 85.9 no effect	257 vs 295 contradicted	181 vs 185 no effect	218 vs 217 no effect
H2 · direction_switches_per_1000 lower is better	0 vs 63.5 supported	13.5 vs 35.5 supported	22.5 vs 1.16e+03 supported	14.5 vs 2.18e+03 supported
H3 · head_on_conflicts lower is better	259 vs 259 no effect	702 vs 1.13e+03 supported	1.49e+03 vs 1.14e+03 contradicted	1.74e+03 vs 1.46e+03 no effect
H4 · mean_service_time lower is better	156 vs 157 no effect	156 vs 136 no effect	144 vs 159 no effect	146 vs 144 no effect
H5 · p95_service_time lower is better	296 vs 286 no effect	297 vs 297 no effect	285 vs 285 no effect	277 vs 277 no effect
H6 · throughput higher is better	0.149 vs 0.149 no effect	0.134 vs 0.0945 supported	0.204 vs 0.152 no effect	0.245 vs 0.217 no effect
	bottleneck	corridors	narrow	medium